

# TIME SERIES MODELING

*Adam Jones, PhD*

*Data Scientist @ Critical Juncture*

---

## TIME SERIES MODELING

---

# LEARNING OBJECTIVES

- Model and predict from time series data using AR, ARMA, or ARIMA models
- Specifically, coding these models in `statsmodels`

---

**COURSE**

---

**PRE-WORK**

---

## PRE-WORK REVIEW

---

- Prior definition and Python functions for moving averages and autocorrelation
- Prior exposure to linear regression with discussion of coefficients and residuals
- `pip install statsmodels` (should be included with Anaconda)

---

**OPENING**

---

# TIME SERIES MODELING

---

# TIME SERIES MODELING

---

- In the last class, we focused on exploring time series data and common statistics for time series analysis
- In this class, we will advance those techniques to show how to predict or forecast forward from time series data
- With a sequence of values (a time series), we will use the techniques in this class to predict a future value

---

# TIME SERIES MODELING

---

- There are many times when you may want to use a series of values to predict a future value
  - The number of sales in a future month
  - Anticipated website traffic when buying a server
  - Financial forecasting
  - The number of visitors to your store during the holidays

## **INTRODUCTION**

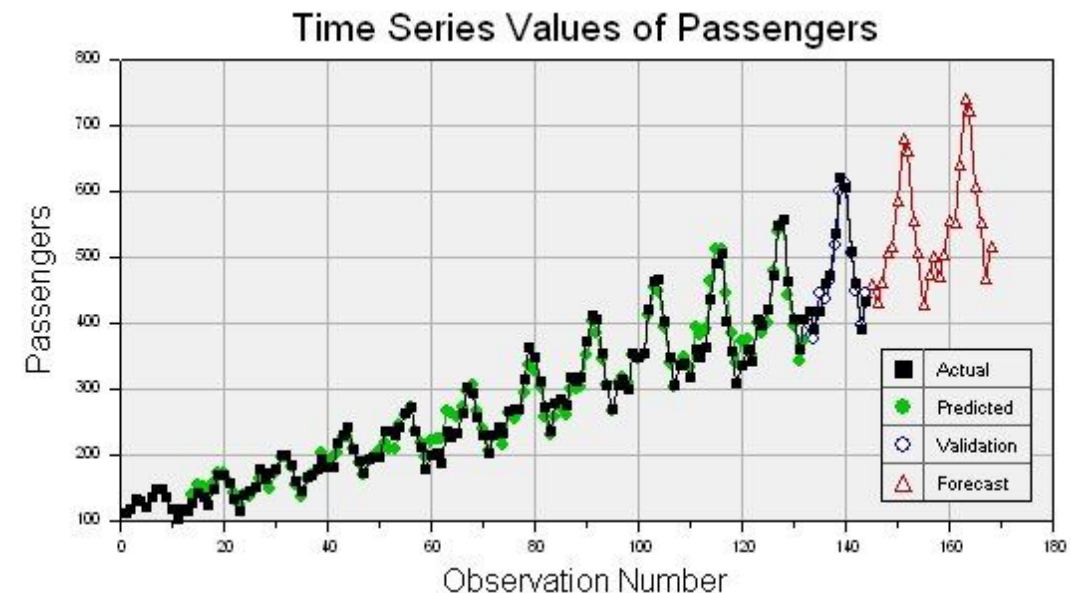
---

# **WHAT ARE TIME SERIES MODELS?**



# WHAT ARE TIME SERIES MODELS?

- ▶ Time series models are models that will be used to predict a future value in the time series
- ▶ **Like** other predictive models, we will use prior history to predict the future
- ▶ **Unlike** previous models, we will use the earlier in time *outcome* variables as *inputs* for predictions



---

## WHAT ARE TIME SERIES MODELS?

---

- **Like** previous modeling exercises, we will have to evaluate the different types of models to ensure we have chosen the best one
- We will want to evaluate on a held-out set or test data to ensure our model performs well on unseen data

---

# WHAT ARE TIME SERIES MODELS?

---

- **Unlike** previous modeling exercises, we won't be able to use standard cross-validation for evaluation
- Since there is a time component to our data, we cannot choose training and test examples at random
- Suppose we did select a random 80% sample of data points for training and a random 20% for testing
  - What could go wrong?

---

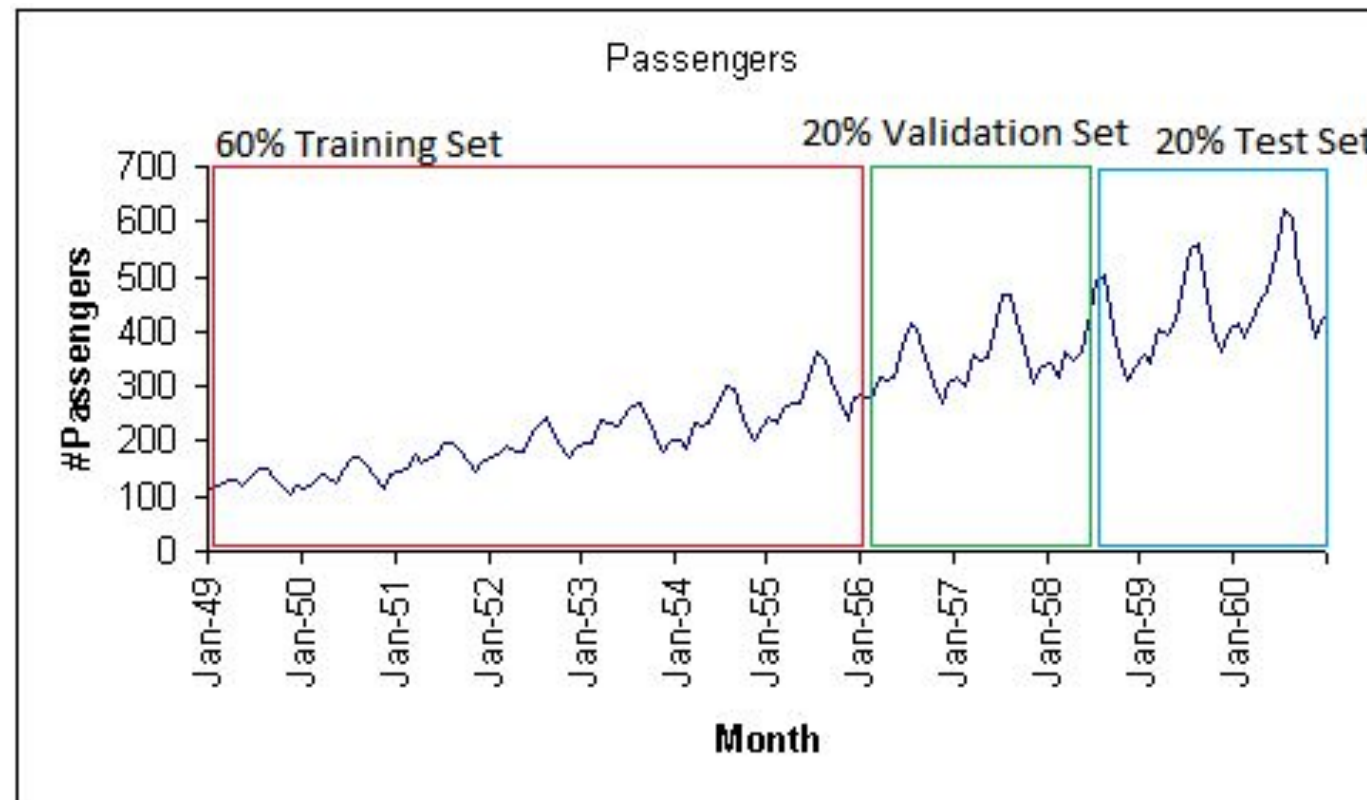
## WHAT ARE TIME SERIES MODELS?

---

- The training dataset would likely contain data from *before* AND *after* a test dataset
- This would not be possible in real life
  - You can't use future, unseen data points when building your model
  - Therefore, it's not a valid test of how our model would perform in practice

# WHAT ARE TIME SERIES MODELS?

- Instead, we will exclusively train on values earlier (in time) in our data and test our model on values at the end of the data period



---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

In our last class, we saw a few statistics for analyzing time series

We looked at *moving averages* to evaluate the local behavior of the time series

1. **Redefine** the moving average and its purpose

### DELIVERABLE

Answers to the above questions

# PROPERTIES FOR TIME-SERIES PREDICTION

- A *moving average* is an average of  $p$  surrounding data points in time

$$F_t = \frac{1}{p} \sum_{k=t-p+1}^{t-p+1} Y_k$$

... to  $t - p + 1$ . This includes  $t, t + 1, t + 2, \dots, t-p, t-p+1$ .

Divide by the  $p$  surrounding data points

Get the  $p$  points from  $t$  ...

---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

We previously looked at auto-correlation to compute the relationship of the data with prior values

1. Recall the definition of autocorrelation and its purpose

### DELIVERABLE

Answers to the above questions



---

# PROPERTIES FOR TIME-SERIES PREDICTION

---

- *Autocorrelation* is how correlated a variable is with itself
- Specifically, how related are variables earlier in time with variables later in time

$$r_k = \frac{\sum_{t=k+1}^n (y_t - \bar{y})(y_{t-k} - \bar{y})}{\sum_{t=1}^n (y_t - \bar{y})^2}$$

- We fix a *lag*,  $k$ , which is how many time points earlier we should use to compute the correlation

---

# PROPERTIES FOR TIME-SERIES PREDICTION

---

- We can use these values to assess how we plan to model our time series
- Typically, for a high quality model, we require some autocorrelation in our data
- We can compute autocorrelation at various lag values to determine how far back in time we need to go

---

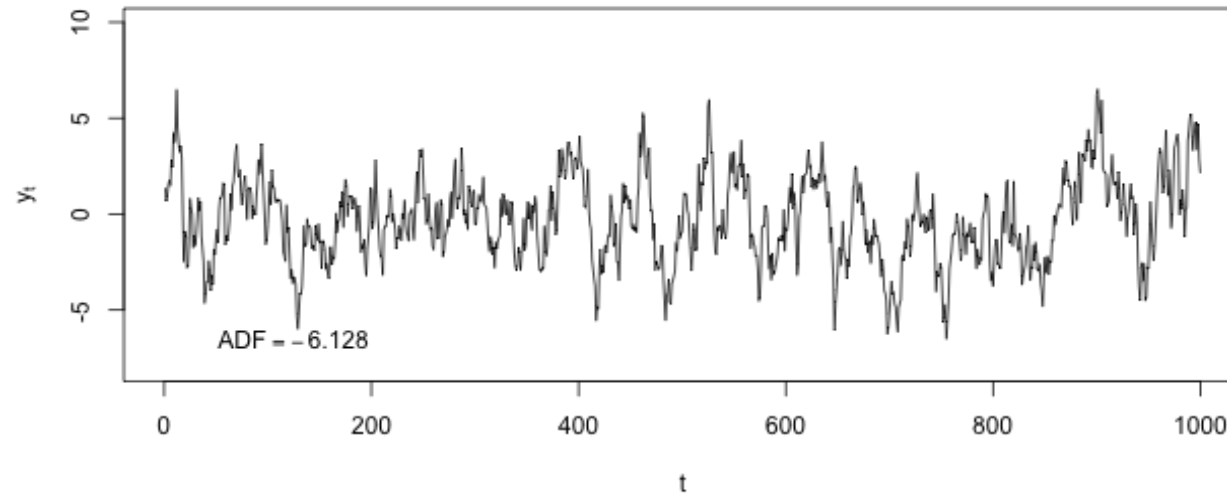
# PROPERTIES FOR TIME-SERIES PREDICTION

---

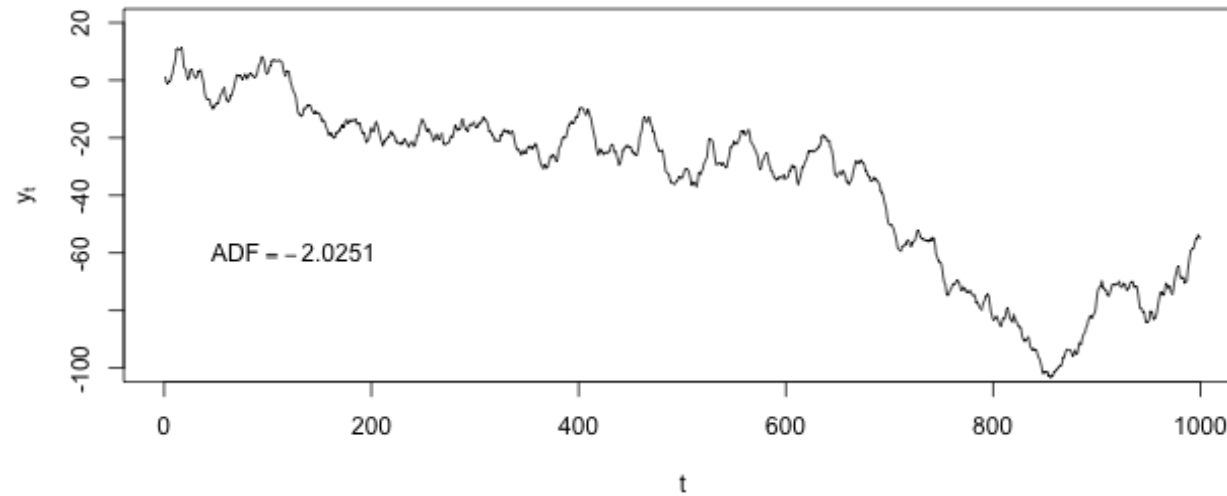
- Many models make an assumption of *stationarity*, assuming the mean and variance of our values is the *same* throughout
- While the values (e.g. of sales) may shift up or down over time, the mean and variance of sales is constant (i.e. there aren't many dramatic swings up or down)
- These assumptions may not represent real world data; we must be aware of that when we are breaking the assumptions of our model

# PROPERTIES FOR TIME-SERIES PREDICTION

**Stationary Time Series**



**Non-stationary Time Series**

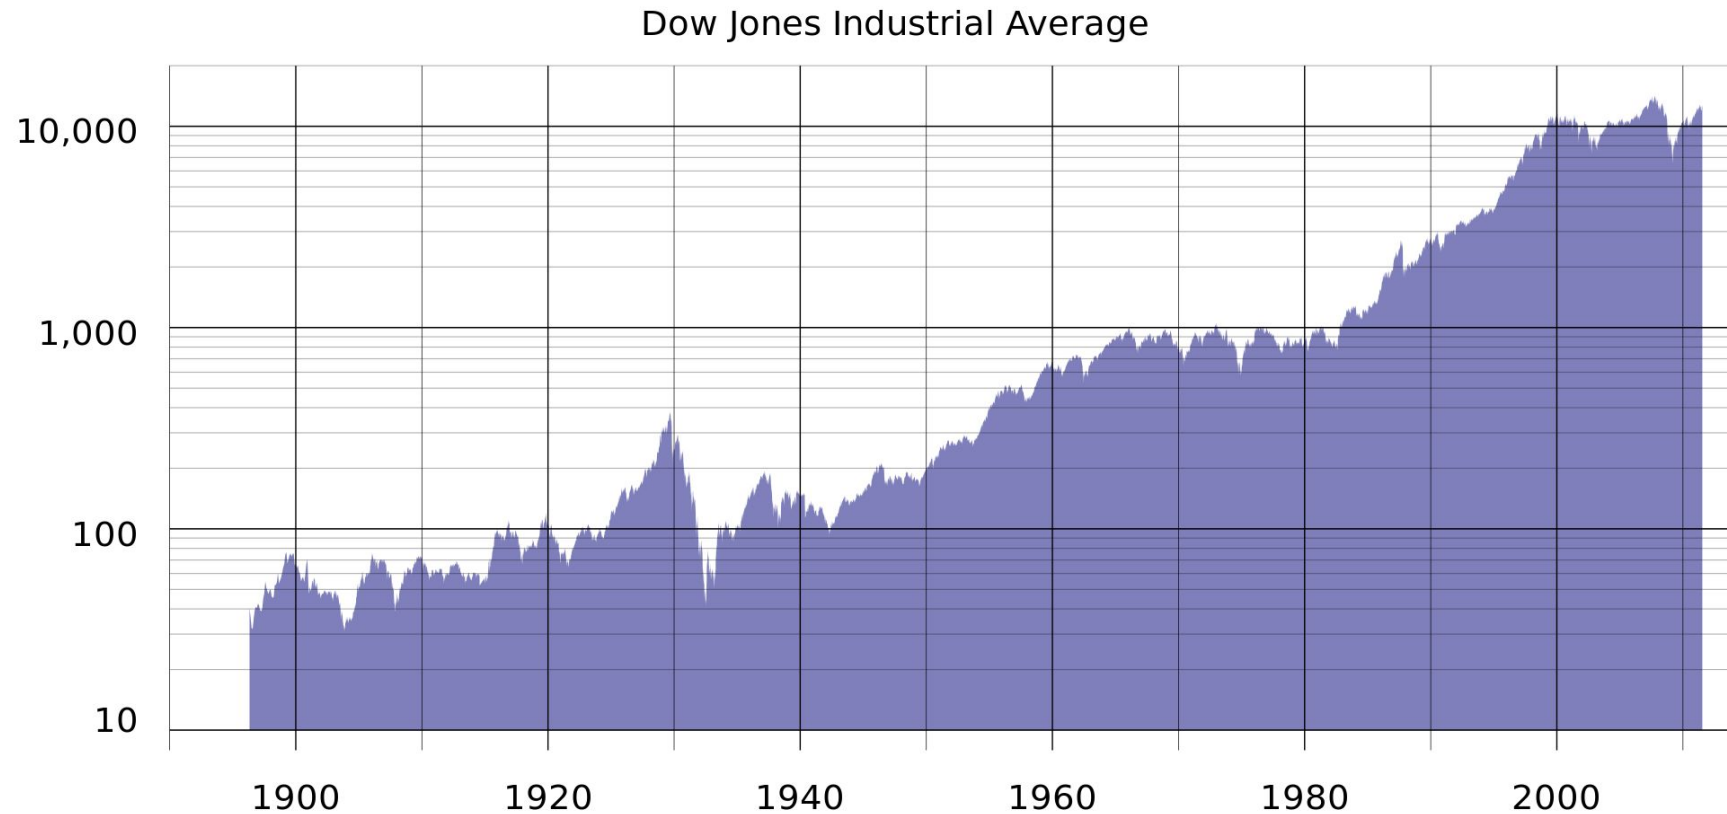


---

# PROPERTIES FOR TIME-SERIES PREDICTION

---

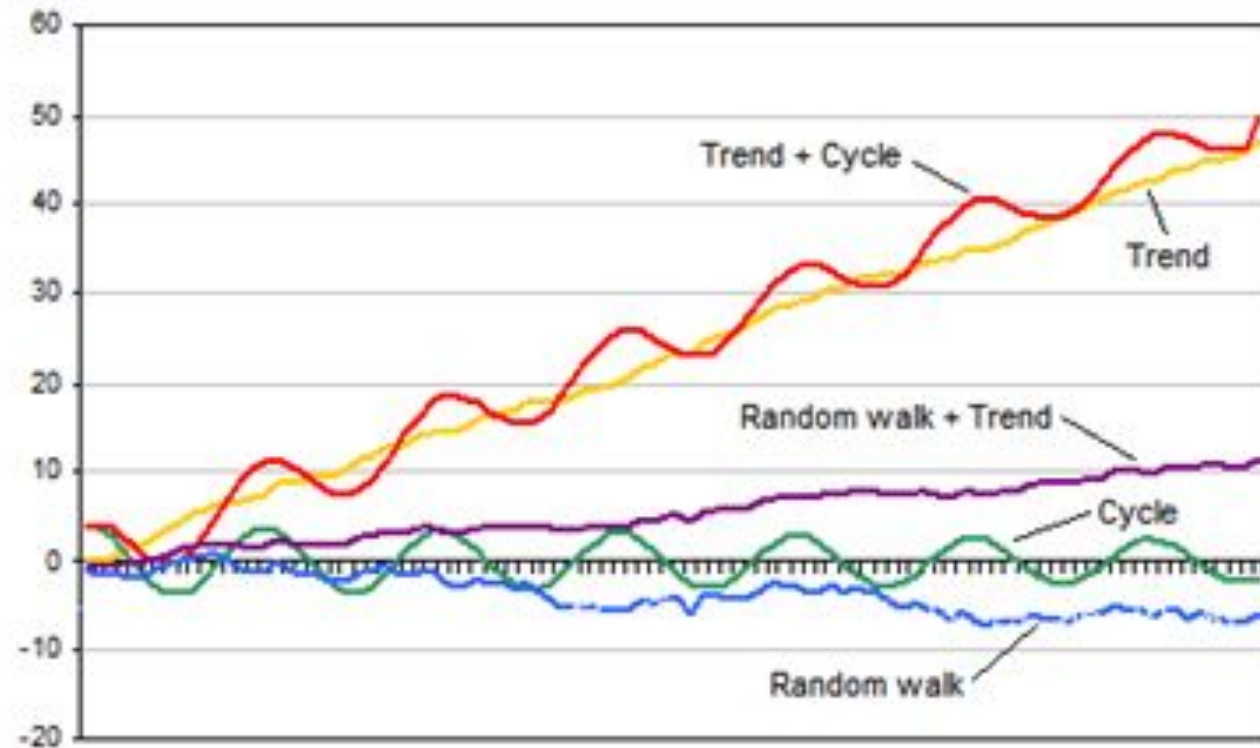
- For example, typical stock or market performance is not stationary
  - In this plot of Dow Jones performance since 1986, the mean is clearly increasing over time



# PROPERTIES FOR TIME-SERIES PREDICTION

- Below are simulated examples of non-stationary time series and why they might occur

**Table 1 Non-stationary behavior**



---

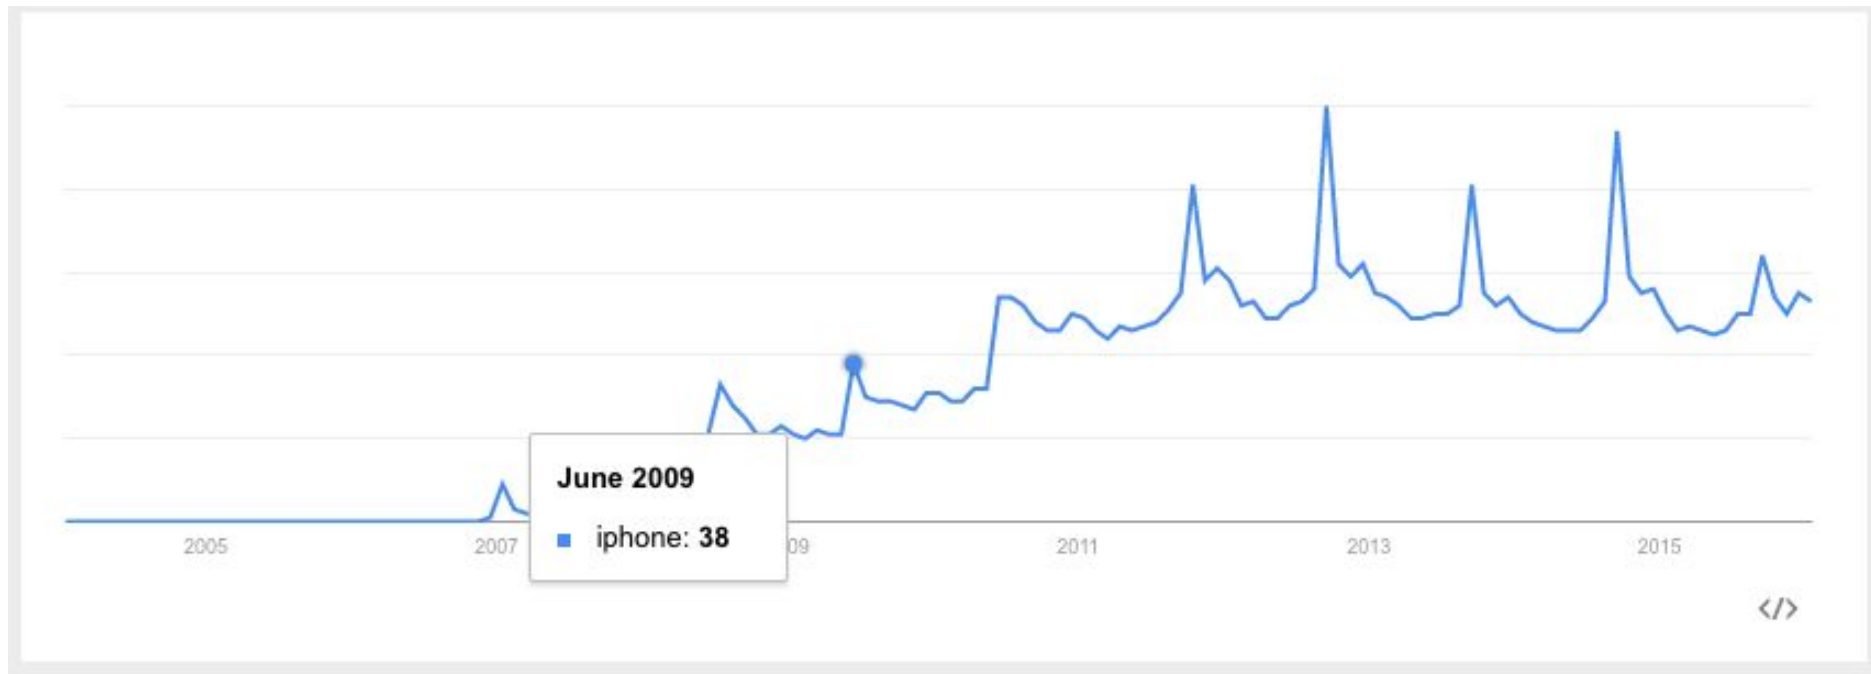
# PROPERTIES FOR TIME-SERIES PREDICTION

---

- Often, if these assumptions don't hold, we can alter our data to make them true
  - Two common methods are *detrending* and *differencing*
- *Detrending* would mean to remove any major trends in our data
- We could do this in many ways, but the simplest is to fit a line to the trend and make a new series that is the difference between the line and the true series

# PROPERTIES FOR TIME-SERIES PREDICTION

- For example, there is a clear upward (non-stationary) trend in google searches for “iphone”

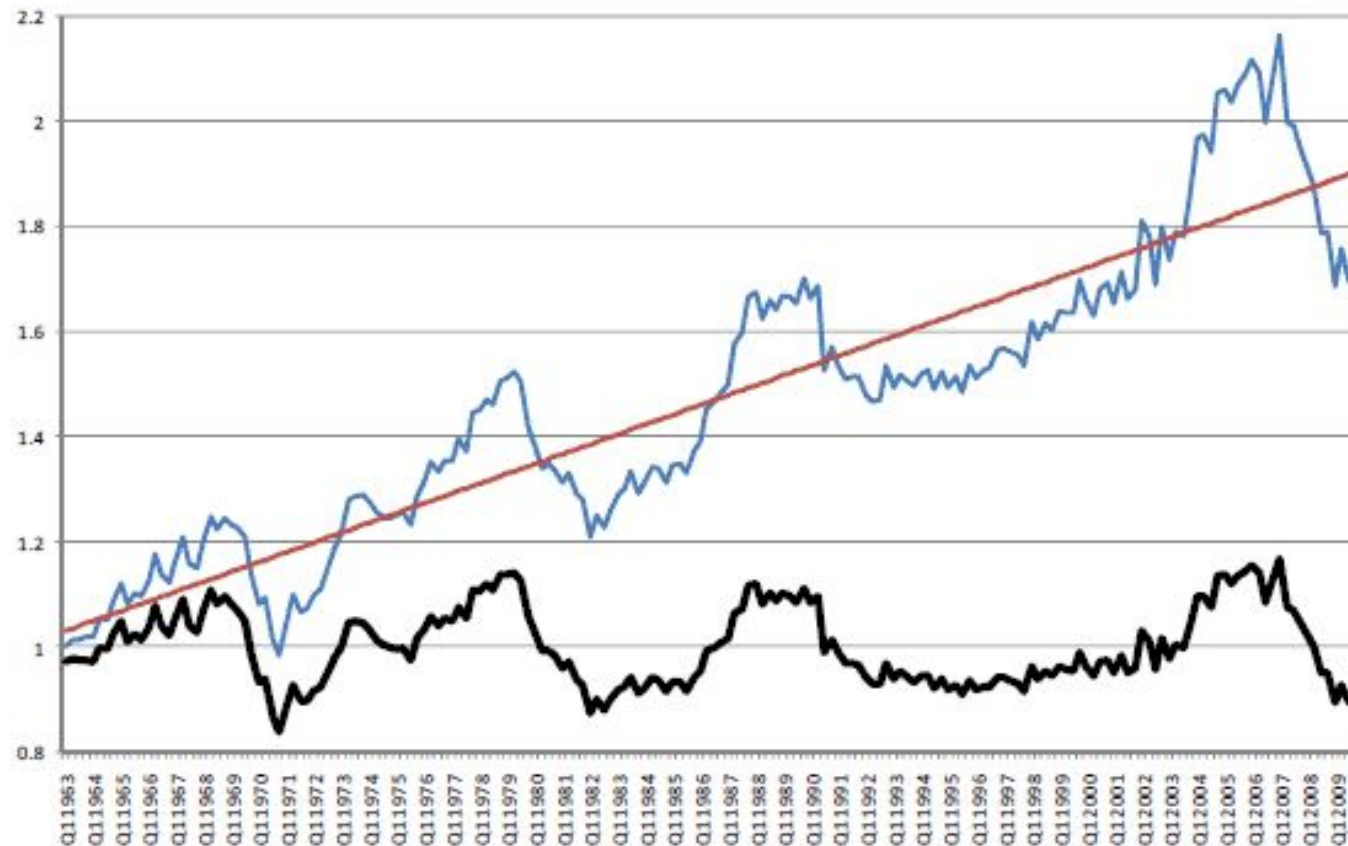


- If we fit a line to this data first, we can create a new series that is the difference between the true number of searches and the predicted searches



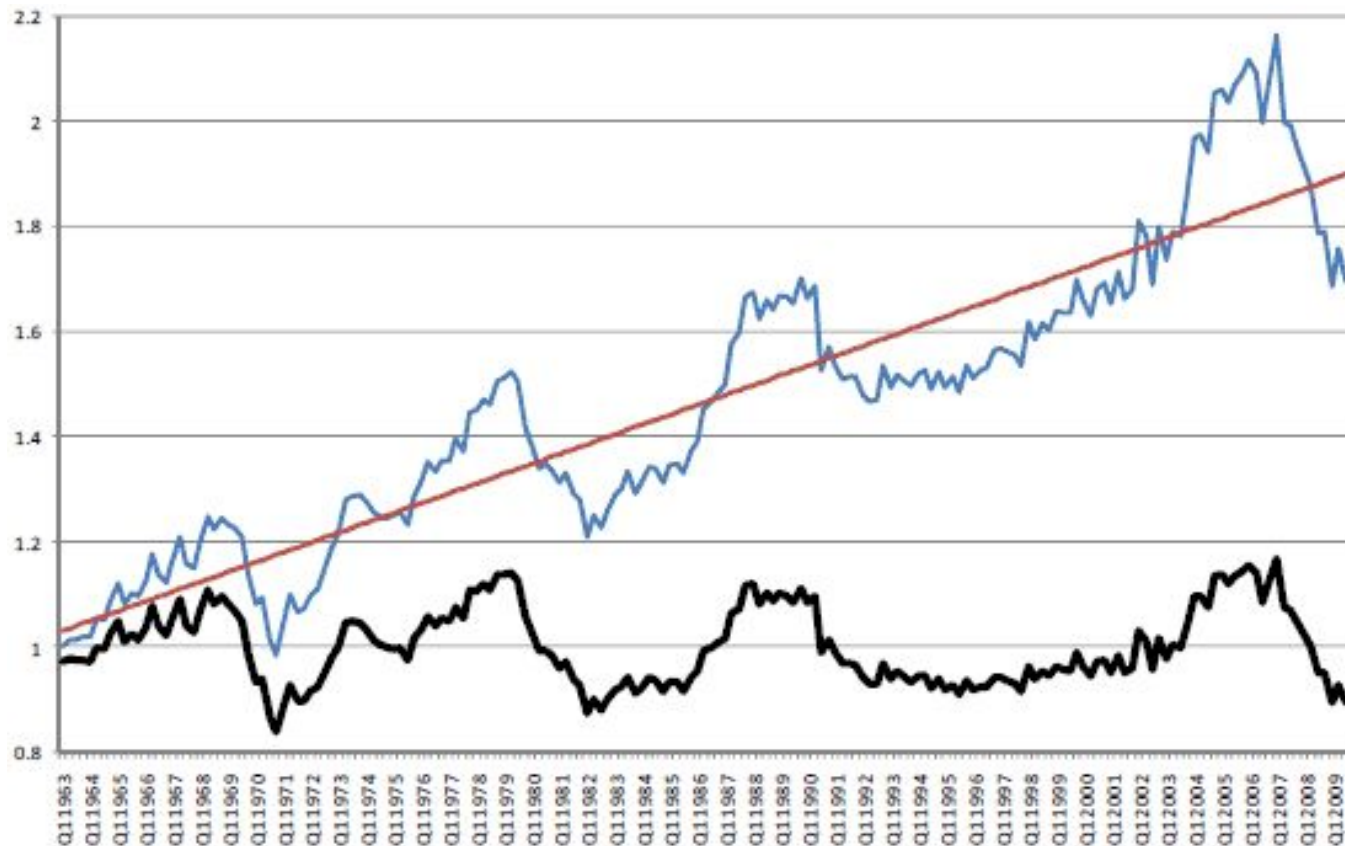
# PROPERTIES FOR TIME-SERIES PREDICTION

- Below is an example where we look at US housing prices over time
- Clearly, there is an upward trend, making the time series non-stationary (ie: the mean house price is increasing)



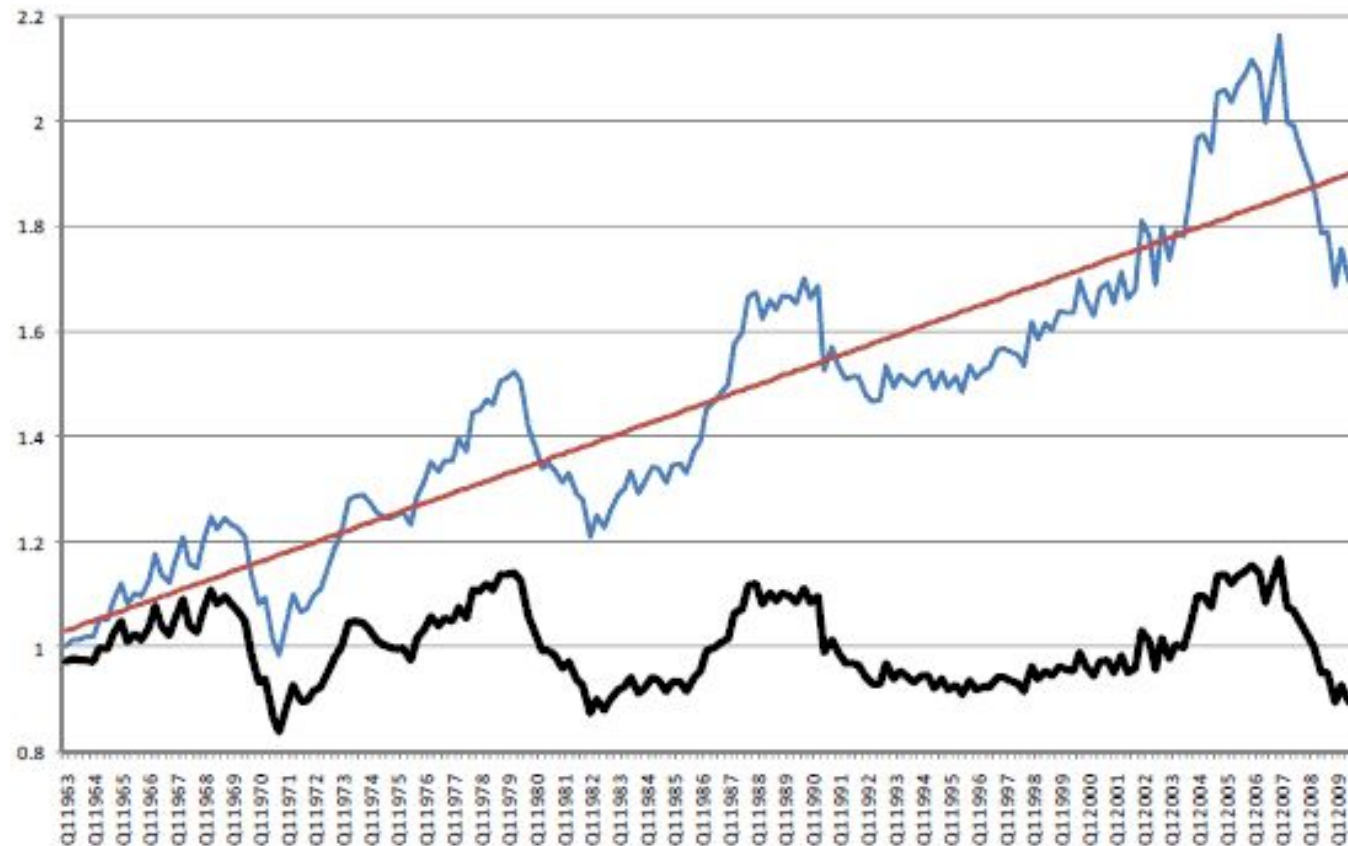
# PROPERTIES FOR TIME-SERIES PREDICTION

- We can fit a line that represents the trend
  - With our trend line, we can subtract the trend line value from the original value to get the bottom figure



# PROPERTIES FOR TIME-SERIES PREDICTION

- The data now has a fixed mean and will be easier to model
- This pattern is similar to mean-scaling our features in earlier models with `StandardScaler`



---

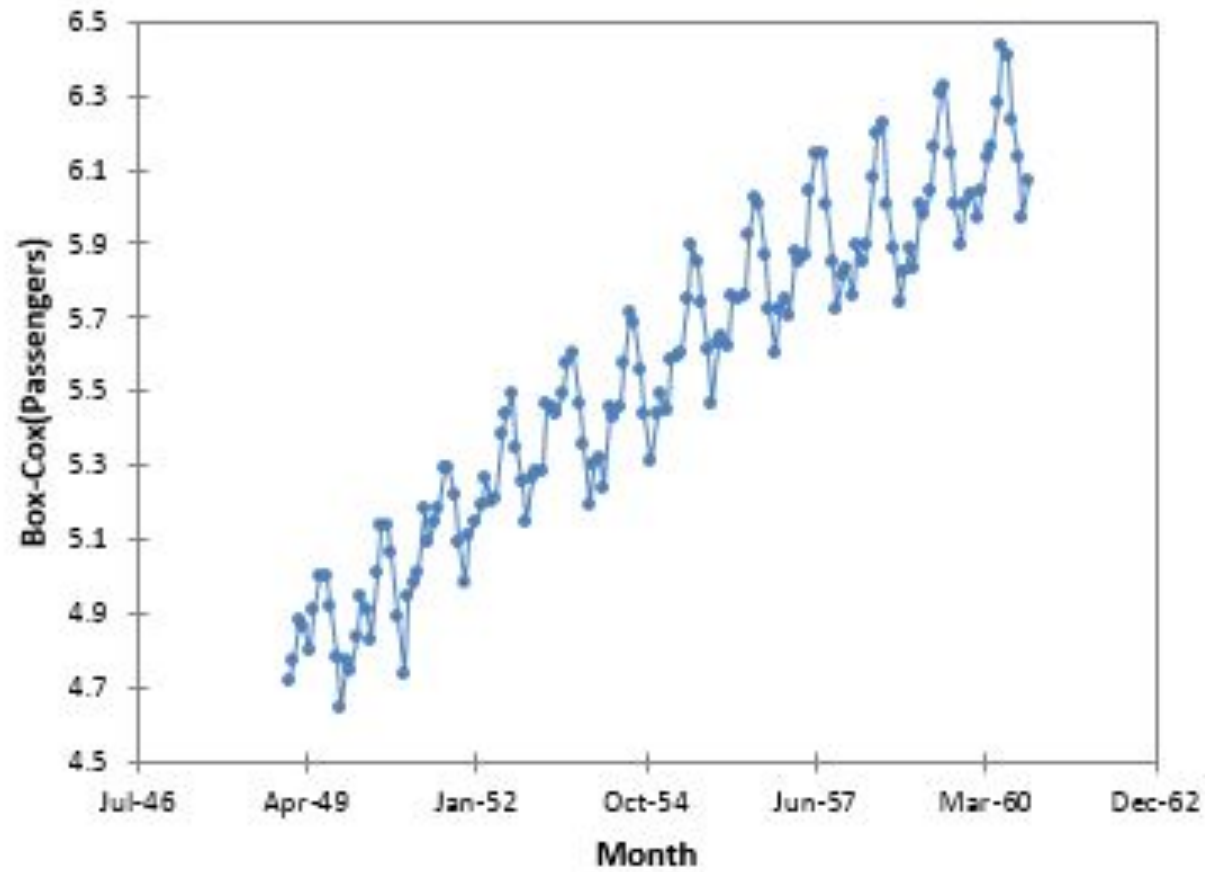
# PROPERTIES FOR TIME-SERIES PREDICTION

---

- A simpler method is *differencing*
  - This is very closely related to the `diff` function we saw in the last class
- Instead of predicting the series (again our non-stationary series), we can predict the difference between two consecutive values

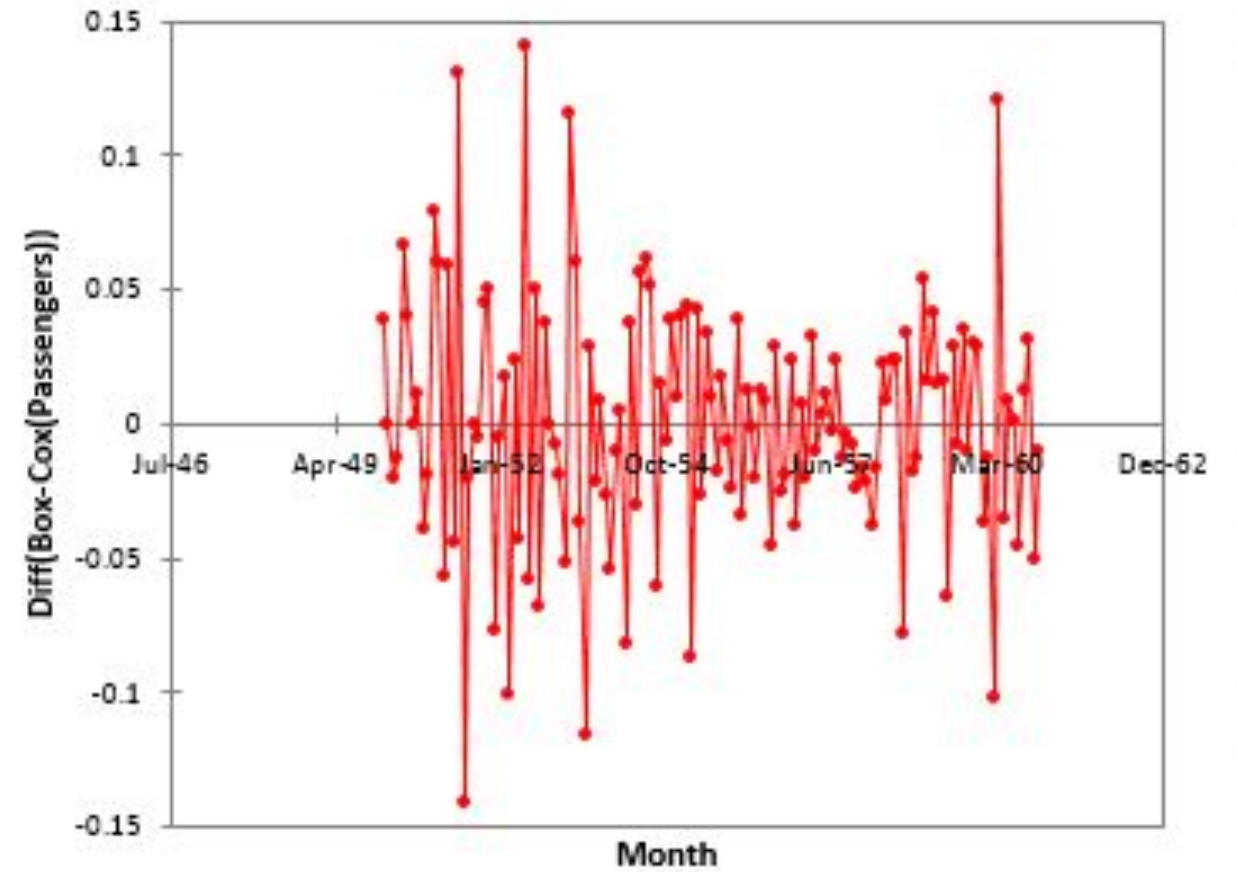
# PROPERTIES FOR TIME-SERIES PREDICTION

Box-Cox(Passengers)



Box-Cox(Passengers)

Differencing (Box-Cox(Passengers))



Box-Cox(Passengers)

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

Non-stationary data is the most common; almost any interesting dataset is non-stationary

1. Can you think of some interesting datasets that might be stationary?

### DELIVERABLE

Answers to the above questions

---

# TIME SERIES MODELS

---

- In the rest of this lesson, we are going to build up to the **ARIMA** time series model
- This model combines the ideas of differencing and two models we will see
  - **AR** - autoregressive models
  - **MA** - moving average models

---

# AR MODELS

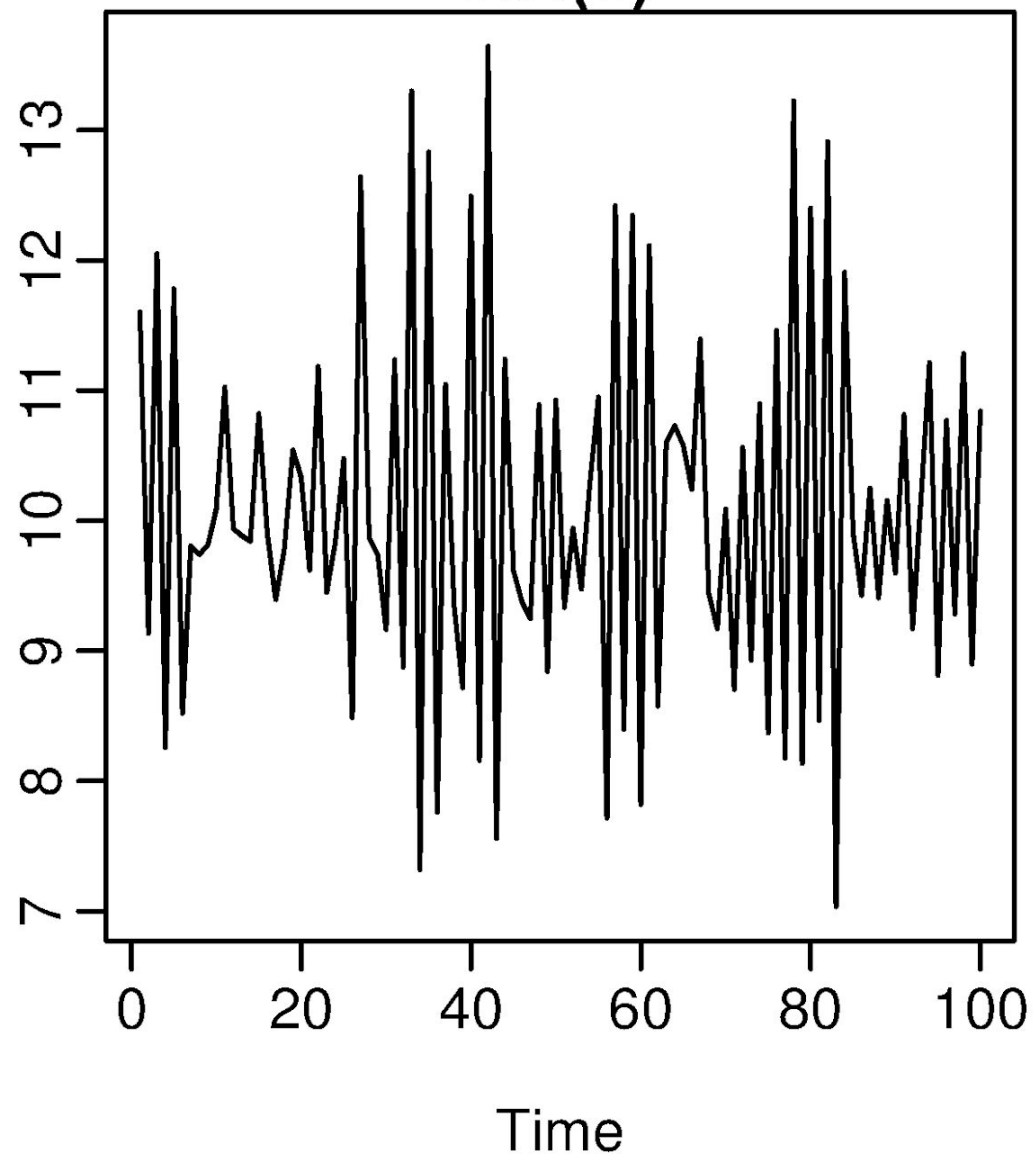
---

- Autoregressive (AR) models are those that use data from previous time points to predict the next
- This is very similar to previous regression models, except as input, we take the previous outcome
- If we are attempting to predict weekly sales, we use the sales from a previous week as input
- Typically, AR models are noted  $AR(p)$  where  $p$  indicates the number of previous time points to incorporate, with  $AR(1)$  being the most common

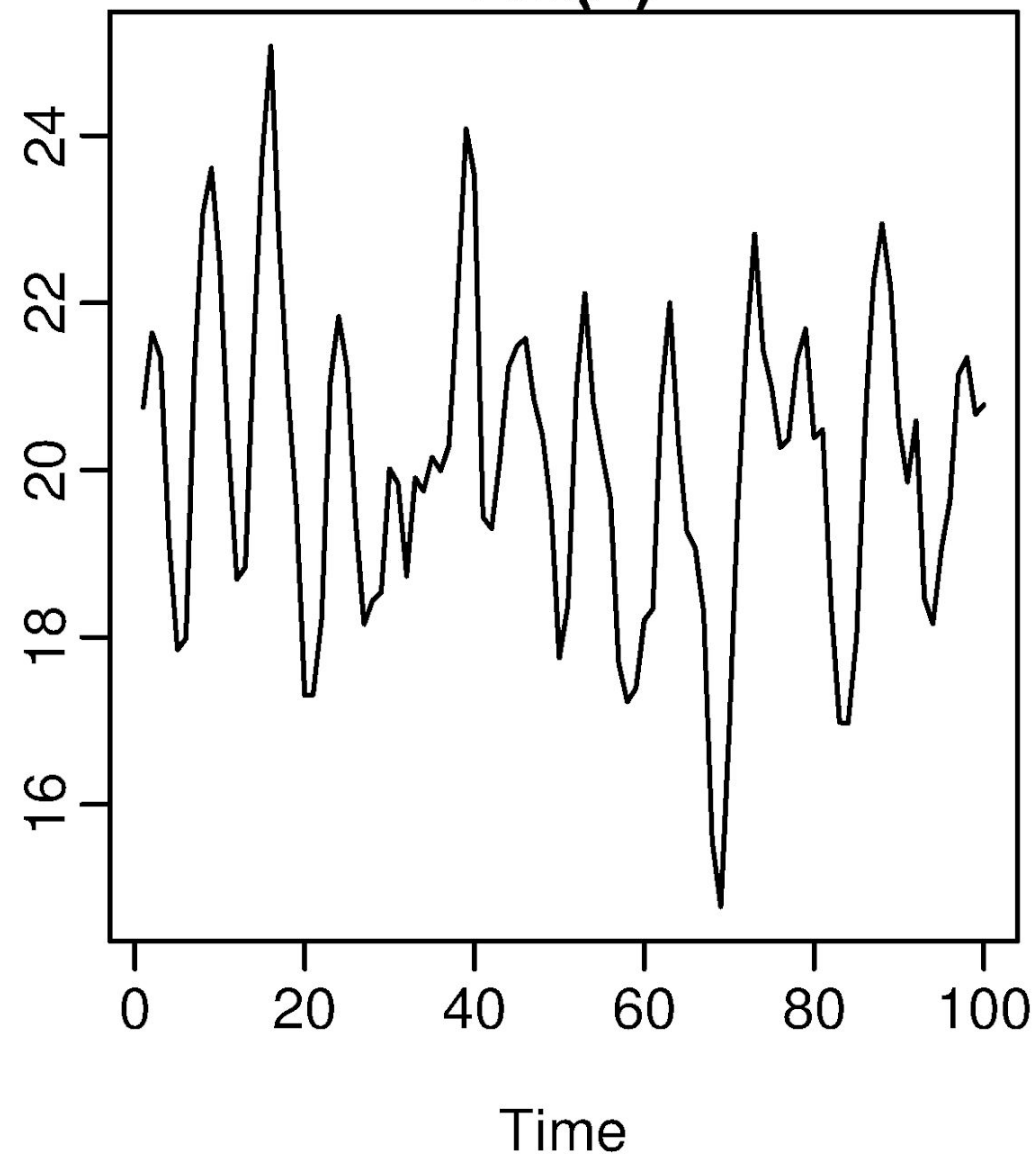


# AR MODELS

AR(1)



AR(2)



---

# AR MODELS

---

- In an autoregressive model, similar to standard regression, we are learning regression coefficients for each of the  $p$  previous values
  - Therefore, we will learn  $p$  coefficients or  $\beta$  values
- If we have a time series of sales per week,  $y_i$ , we can regress each  $y_i$  from the last  $p$  values

$$y_i = \beta_0 + \beta_1 y_{i-1} + \beta_2 y_{i-2} + \dots + \beta_p y_{i-p} + \varepsilon$$

- As with standard regression, our model assumes that each outcome variable is a linear combination of the inputs and a random error term

---

# AR MODELS

---

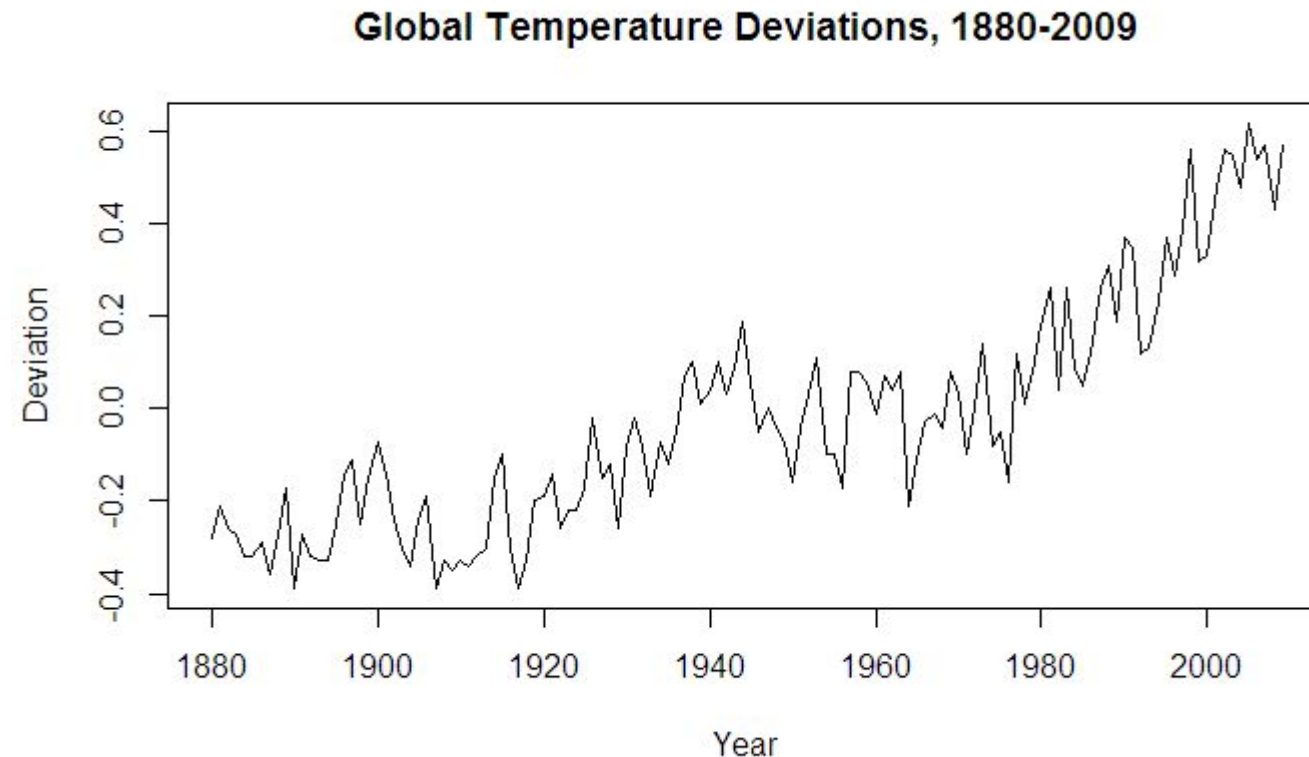
- For an AR(1) model, we will learn a single coefficient
- This coefficient,  $\beta$ , will tell us the relationship between the previous value,  $Y_{t-1}$ , and the next value,  $Y_t$

$$Y_t = \beta \cdot Y_{t-1}$$

# AR MODELS

---

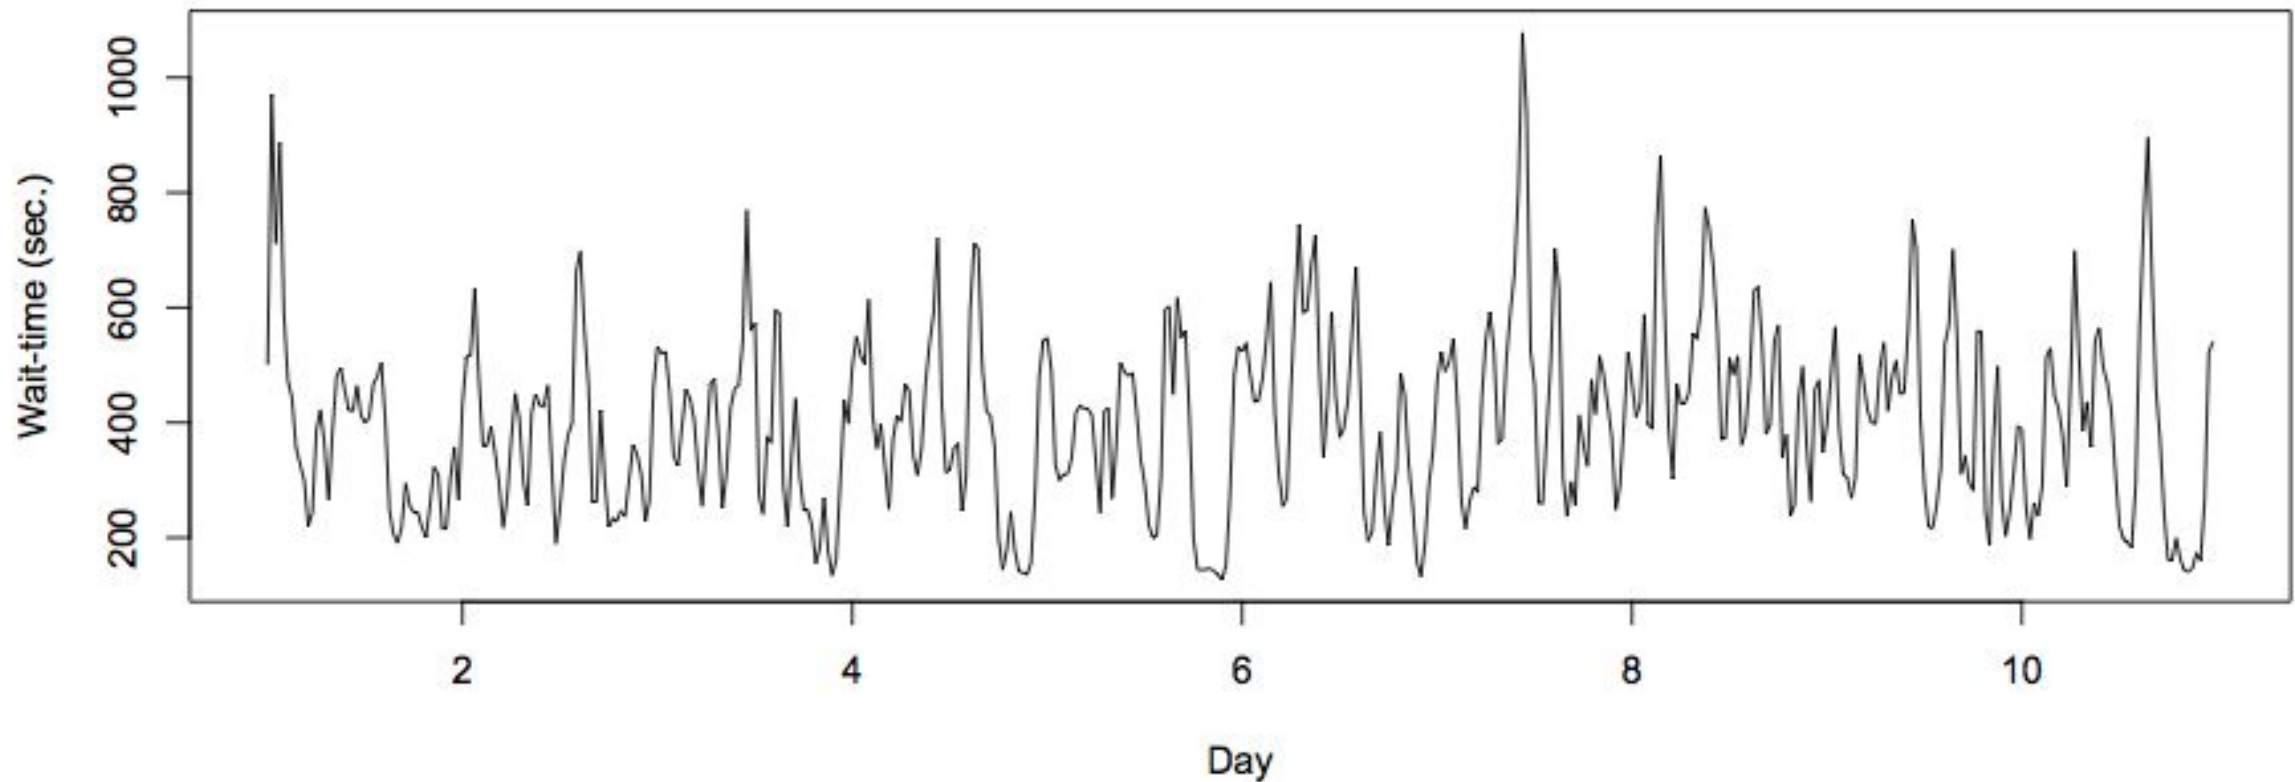
- A value  $> 1$  would indicate a growth over previous values
- This would typically represent non-stationary data, since if we compound the increases, the values are continually increasing



# AR MODELS

---

- Values between 1 and -1 represent increasing and decreasing patterns from previous patterns



---

## AR MODELS

---

- As with other models, interpretation of the model becomes more complex as we add more factors
- Going from AR(1) to AR(2) can add significant *multicollinearity*

---

## AR MODELS

---

- Recall that *autocorrelation* is the correlation of a value with its series *lagged* behind
- A model with high correlation implies that the data is highly dependent on previous values and an autoregressive model would perform well

---

## AR MODELS

---

- Autoregressive models are useful for learning falls or rises in our series
- This will weight together the last few values to make a future prediction
- Typically, this model type is useful for small-scale trends such as an increase in demand or change in tastes that will gradually increase or decrease the series



---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

1. If we observe an autocorrelation near 1 for lag 1, what do we expect the single coefficient in an AR(1) model to be?  $>1$ , between 0 and 1, or  $<1$ ?
2. What if we observe an autocorrelation of 0?

### DELIVERABLE

Answers to the above questions

---

## MA MODELS

---

- **Moving average (MA) models**, as opposed to AR models, do not take the previous outputs (or values) as inputs
  - They take the previous error terms
- We will attempt to predict the next value based on the overall average and how far off our previous predictions were

---

## MA MODELS

---

- This model is useful for handling specific or abrupt changes in a system
- AR models slowly incorporate changes in the system by combining previous values; MA models use prior errors to quickly incorporate changes
- This is useful for modeling a sudden occurrence - something going out of stock or a sudden rise in popularity affecting sales

---

## MA MODELS

---

- As in AR models, we have an order term,  $q$ , and we refer to our model as MA( $q$ )
  - The moving average model is dependent on the last  $q$  errors
- If we have a time series of sales per week,  $y_i$ , we can regress each  $y_i$  from the last  $q$  error terms

$$y_i = \text{mean} + \beta_1 \varepsilon_{i-1} + \beta_2 \varepsilon_{i-2} + \dots + \beta_q \varepsilon_{i-q}$$

- We include the mean of the time series (that's why it's called a moving average) as we assume the model takes the mean value of the series and randomly jumps around it

---

## MA MODELS

---

- Of course, we don't have error terms when we start - where do they come from?
- This requires a more complex fitting procedure than we have seen previously
- We need to iteratively fit a model (perhaps with random error terms), compute the errors and then refit, again and again

---

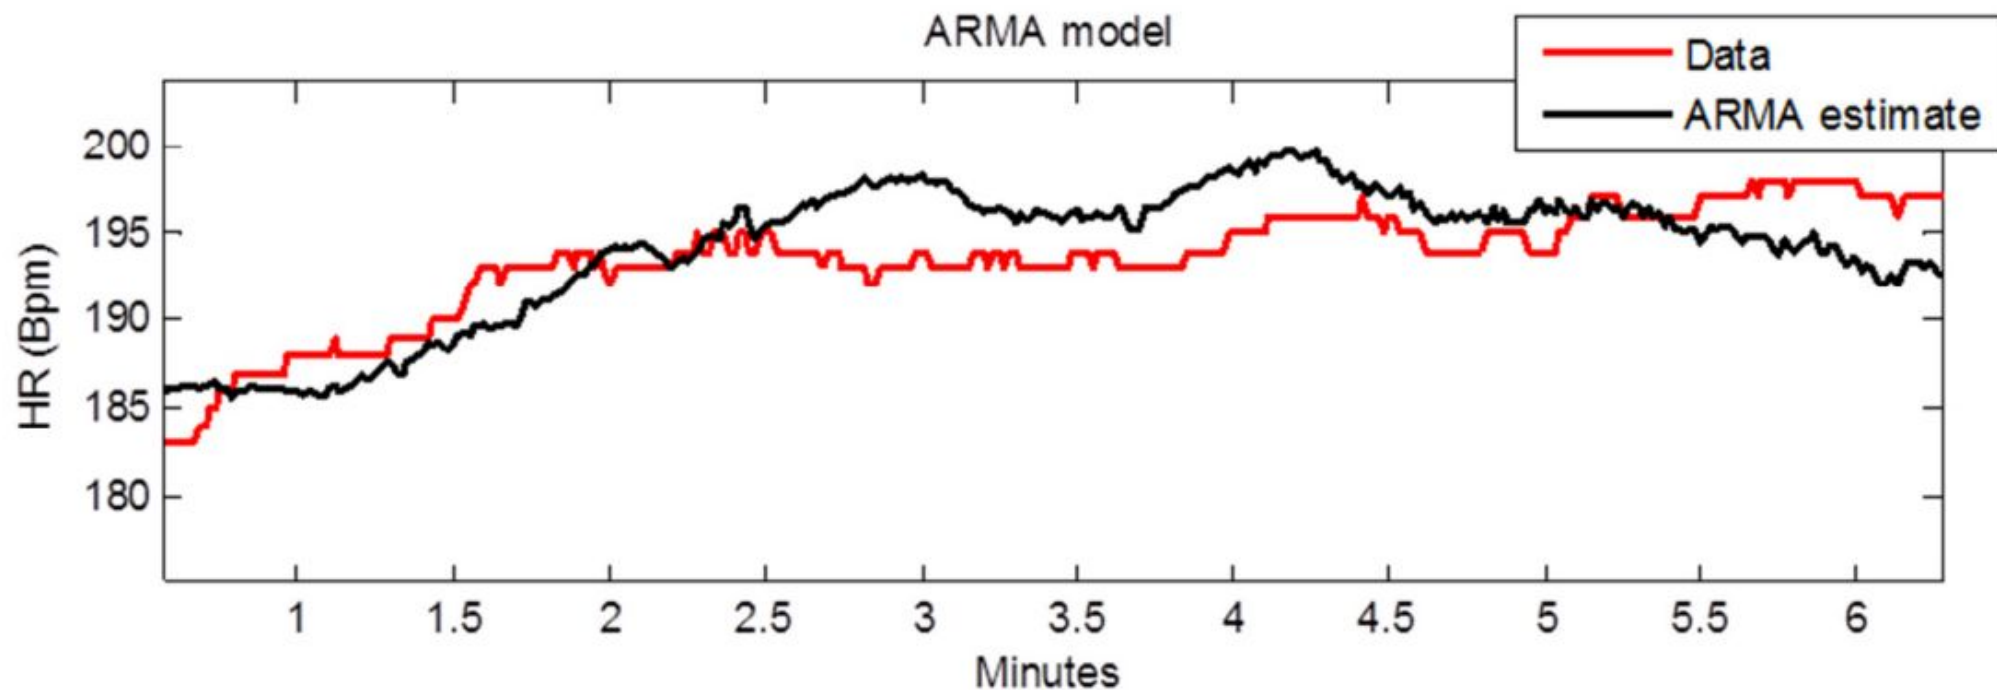
## MA MODELS

---

- In this model, we learn  $q$  coefficients
- In an MA(1) model, we learn one coefficient
- This value indicates the impact of how our previous error term on the next prediction

# ARMA MODELS

- **ARMA** (pronounced 'R-mah') models combine the autoregressive and moving average models
- An ARMA(p,q) model is simply a combination (sum) of an AR(p) model and MA(q) model



---

# ARMA MODELS

---

- We specify two model settings,  $p$  and  $q$ , which correspond to combining an AR( $p$ ) model with an MA( $q$ ) model
- Incorporating both models allows us to mix two types of effects
  - AR models slowly incorporate changes in preferences, tastes, and patterns
  - Moving average models base their prediction on the prior error, allowing to correct sudden changes based on random events
    - supply, popularity spikes, etc.



---

# ARIMA MODELS

---

- **ARIMA** (pronounced ‘uh-ri-mah’) is an **AutoRegressive Integrated Moving Average** model
- In this model, we learn an ARMA(p,q) model to predict *the difference* of the series (as opposed to the value of the series)

---

# ARIMA MODELS

---

- Recall the pandas `diff` function
  - This computes the difference between two consecutive values
- In an ARIMA model, we attempt to predict this difference instead of the actual values

$$y_t - y_{t-1} = \text{ARIMA}(p,q)$$

- This handles the stationarity assumption we wanted for our data
- Instead of detrending or differencing manually, the model does this

---

# ARIMA MODELS

---

- An ARIMA model has three parameters and is specified ARIMA( $p$ ,  $d$ ,  $q$ )
  - $p$  is the order of the autoregressive component
  - $q$  is the order of the moving average component
  - $d$  is the degree of differencing
- $d$  was 1 in our prior example. For  $d=2$ , our model would be

$$\text{diff}(\text{diff}(y)) = (y_t - y_{t-1}) - (y_{t-1} - y_{t-2}) = \text{ARIMA}(p, q)$$

---

# ARIMA MODELS

---

- Compared to an ARMA model, ARIMA models do **not** rely on the underlying series being stationary
- The differencing operation can *convert* the series to one that is stationary
- Instead of attempting to predict values over time, our new series is the difference in values over time
- Since ARIMA models *include differencing*,
  - they can be used on a broader set of data without the assumption of a constant mean

**DEMO**

---

# TIME SERIES MODELING IN STATSMODELS

---

# TIME SERIES MODELING IN STATSMODELS

---

- To explore time series models, we will continue to use the Rossmann sales data
- This dataset has sales data for every Rossmann store for a 3-year period and indicators for holidays and basic store information

---

# TIME SERIES MODELING IN STATSMODELS

---

- In the last class, we saw that we could plot the sales data at a particular store to identify how the sales changed over time
- We also computed autocorrelation for the data at varying lag periods
- This helps us identify if previous timepoints are predictive of future data and which time points are most important - the previous day, week, or month

---

# TIME SERIES MODELING IN STATSMODELS

---

```
import pandas as pd
```

```
# Load the data and set the DateTime index  
data = pd.read_csv('./datasets/rossmann.csv', skipinitialspace=True)
```

```
data['Date'] = pd.to_datetime(data['Date'])  
data.set_index('Date', inplace=True)
```

```
# Filter to Store 1  
store1_data = data[data.Store == 1]
```

```
# Filter to open days  
store1_open_data = store1_data[store1_data.Open==1]
```

```
# Plot the sales over time  
store1_open_data[['Sales']].plot()
```



---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

1. Compute the autocorrelation of Sales in Store 1 for lag 1 and 2.
2. Will we be able to use a predictive model, particularly an autoregressive one?

### DELIVERABLE

Answers to the above questions

---

# TIME SERIES MODELING IN STATSMODELS

---

```
store1_data.Sales.autocorr(lag=1) # -0.12  
store1_data.Sales.autocorr(lag=2) # -0.03
```

- We do see some minimal correlation in time, implying an AR model can be useful
  - An easier way to diagnose this may be to plot many autocorrelations at once

---

# TIME SERIES MODELING IN STATSMODELS

---

```
from pandas.tools.plotting import autocorrelation_plot
%matplotlib inline

autocorrelation_plot(store1_data.Sales)
```

- This shows a typical pattern of an autocorrelation plot, that it should decrease to 0 as lag increases
  - However, it's hard to observe exactly what the values are

---

# TIME SERIES MODELING IN STATSMODELS

---

- In this class, we will use `statsmodels` to code AR, MA, ARMA, and ARIMA models
- `statsmodels` provides a nice summary utility to help us diagnose models

---

# TIME SERIES MODELING IN STATSMODELS

---

- statsmodels also has a better autocorrelation plot that allows us to look at fixed number of lag values

```
from statsmodels.graphics.tsaplots import plot_acf  
plot_acf(store1_data.Sales, lags=10)
```

- Here we observe autocorrelation at 10 lag values
  - 1 and 2 are what we saw before

---

# TIME SERIES MODELING IN STATSMODELS

---

- We also see a larger spike at 7 (the seventh day in the week)
- If we observed a handful of random distributed spikes, a moving average model would be useful

```
plot_acf(store1_data.Sales, lags=25)
```

- We can expand the window to 25 days to see that the random spikes occur regularly at 7 days
  - What does this mean?

---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- To explore AR, MA, and ARMA models, we will use `sm.tsa.ARMA`
- Remember, an ARMA model is a combination of autoregressive and moving average models
- We can train an AR model by turning off the MA component ( $q=0$ )

```
from statsmodels.tsa.arima_model import ARMA
```

```
store1_sales_data = store1_open_data[['Sales']].astype(float)  
model = ARMA(store1_sales_data, (1, 0)).fit()  
model.summary()
```

---

## AR, MA, AND ARMA MODELS IN STATSMODELS

---

- By passing `(1, 0)` in the second argument, we are fitting an ARMA model with  $p=1$ ,  $q=0$  (this is the same as an AR(1) model)
- In this AR(1) model, we learn an intercept (or base sales) value
- Additionally, we learn a coefficient that tells us how to include the latest sales value
- In this case, we add an intercept of  $\sim 4700$  to  $0.68$  times the previous month's sales
  - Note that the coefficient is not equal to the lag 1 autocorrelation
  - This implies the data is **not** stationary



---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- We can learn an AR(2) model, which regresses each sales value on the last two

```
model = ARMA(store1_sales_data, (2, 0)).fit()  
model.summary()
```

- In this case, we learn two coefficients, which tell us the effect of the last two sales values on the current sales
- While this model may perform better, it may be more difficult to interpret

---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

To start to diagnose the model, we want to look at residuals.

1. What are residuals?
2. In linear regression, what did we expect of residuals?

### DELIVERABLE

Answers to the above questions

---

## AR, MA, AND ARMA MODELS IN STATSMODELS

---

- Residuals are the errors of the model or how off our predictions are
- Ideally, we want randomly distributed errors that are small
- If the errors are large, our model does not perform well
- If the errors have a pattern, particularly over time, we may have overlooked something in the model or have periods of time that are different than the rest of the dataset

---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- We can use `statsmodels` to plot the residuals

```
model.resid.plot()
```

- Here we see large spikes at the end of each year, indicating that our model does not account for the holiday spikes
- Our model considers a short period of time, so it does not take into account the longer seasonal pattern

---

# AR, MA, AND ARMA MODELS IN STATSMODELS

---

- We can also plot the autocorrelations of the residuals
  - In an ideal world, these would all be near 0 and appear random

```
plot_acf(model.resid, lags=50)
```

- This plot shows a problem: the errors are increasing and decreasing every week in a clear pattern
- We may need to expand our model

---

## AR, MA, AND ARMA MODELS IN STATSMODELS

---

- To expand this AR model to an ARMA model, we can include the moving average component as well

```
model = ARMA(store1_sales_data, (1, 1)).fit()  
model.summary()
```

- Now we learn two coefficients, one for the AR(1) component and one for the MA(1) component

---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

1. Take a moment to look at the coefficients of our new model
2. Offer an interpretation of this model

### DELIVERABLE

Answers to the above questions

---

## AR, MA, AND ARMA MODELS IN STATSMODELS

---

- Remember that this is an AR(1) + MA(1) model
  - The AR coefficient represents dependency on the last value and the MA component represents any spikes independent of the last value
- The coefficients here are 0.69 for the AR component and -0.03 for the MA component
- The AR coefficient is the same as before (decreasing values)
- The MA component is fairly small (which we should have expected from the autocorrelation plots)



---

# ARIMA MODELS IN STATSMODELS

---

- We can also use `statsmodels` to fit ARIMA models
  - Let's start by using `ARIMA(1, 0, 1)` to fit an `ARMA(1, 1)` model

```
from statsmodels.tsa.arima_model import ARIMA
```

```
model = ARIMA(store1_sales_data, (1, 0, 1)).fit()  
model.summary()
```

- We can see that this model is the same as our previous ARMA model

---

# ARIMA MODELS IN STATSMODELS

---

- We can also fit a true ARIMA model to predict the difference of the series

```
model = ARIMA(store1_sales_data, (1, 1, 1)).fit()  
model.summary()
```

- We can remove the MA component since it does not appear to be useful

```
model = ARIMA(store1_sales_data, (1, 1, 0)).fit()  
model.summary()
```

- We now have an AR(1) model on the differenced series with a coefficient of -0.18

---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

1. Does this model match the lag 1 autocorrelation of the differenced series?
2. Is the data stationary?

### DELIVERABLE

Answers to the above questions

---

# ARIMA MODELS IN STATSMODELS

---

- We can compute the lag 1 autocorrelation of the differenced series and see if they match

```
store1_sales_data.Sales.diff(1).autocorr(1) #-0.181
```

- We can also plot it to see the difference

```
store1_sales_data.Sales.diff(1).plot()
```

- They match!
  - Note that this is generally true, but the variance is NOT constant
  - It's mostly the same throughout the series except around the holidays

---

# ARIMA MODELS IN STATSMODELS

---

- With our models, we can also plot our predictions against the true series using the `plot_predict` function
- We can compare the last 50 days of true values against our predictions

```
model.plot_predict(0, 50)
```

- The function takes two arguments, the start and end index of the dataframe to plot
  - Here, we are plotting the last 50 values

---

# ARIMA MODELS IN STATSMODELS

---

- To plot earlier values with our predictions continuing where the true values stop, we can do the following

```
import matplotlib.pyplot as plt
```

```
fig, ax = plt.subplots()  
ax = store1_sales_data['2014'].plot(ax=ax)
```

```
fig = model.plot_predict(0, 200, ax=ax, plot_insample=False)
```

- This plots true values in 2014 and our predictions 200 days out from 2014

---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS

We can revisit our diagnostics to check that our models are working well.

1. Plot the **residuals** and **autocorrelation of the residuals**
2. Are there *patterns* or *outliers*?

### DELIVERABLE

Answers to the above questions

---

# ARIMA MODELS IN STATSMODELS

---

- The two previous problems remain:
  - Large errors around the holiday period
  - Errors with high autocorrelation



---

# ARIMA MODELS IN STATSMODELS

---

- We can adjust the AR component of the model to adjust for a piece of this
  - Let's increase the lag to 7

```
model = ARIMA(store1_sales_data, (7, 1, 2)).fit()  
model.summary()
```

```
plot_acf(model.resid, lags=50)
```

- This removes some of the autocorrelation in the residuals but large discrepancies still exist
- However, they exist where we are breaking our model assumptions

---

# ACTIVITY: KNOWLEDGE CHECK

---

## ANSWER THE FOLLOWING QUESTIONS



### EXERCISE

1. Alter the time period of predictions and the  $p$ ,  $d$ , and  $q$  parameters.
2. Do any of these improve diagnostics?
3. What does changing  $p$  and  $q$  imply based upon the autocorrelation plot?
4. How about changing  $d$ ?

## DELIVERABLE

Answers to the above questions

---

# ARIMA MODELS IN STATSMODELS

---

- Increasing  $p$  increases the dependency on previous values further (longer lag)
  - But our autocorrelation plots show this isn't necessary past a certain point
- Increasing  $q$  increases the likelihood of an unexpected jump at a handful of points
  - The autocorrelation plots show this doesn't help past a certain point
- Increasing  $d$  increases differencing, but  $d=1$  moves our data towards stationarity (other than a few points)
  - $d=2$  would imply an exponential trend which we don't have here

---

# ARIMA MODELS IN STATSMODELS

---

- There are variants of ARIMA that will better handle the seasonal aspect of our data
  - This is referred to as Seasonal ARIMA
- These models fit **two** ARIMA models, one on the current frequency (daily in our example) and another on the seasonal frequency (maybe monthly or yearly patterns)
- Additionally, issues with seasonality could be handled by preprocessing tricks such as detrending

## **INDEPENDENT PRACTICE**

---

# **WALMART SALES DATA**

# ACTIVITY: WALMART SALES DATA



## EXERCISE

### DIRECTIONS (50 minutes)

We will analyze the weekly sales data from Walmart over a two year period from 2010 to 2012. The data is separated by store and department, but we will focus on analyzing one store for simplicity.

To read in the data

```
import pandas as pd
import numpy as np

%matplotlib inline

data = pd.read_csv('./datasets/train.csv')
data.set_index('Date', inplace=True)
data.head()
```

# ACTIVITY: WALMART SALES DATA

---

## DIRECTIONS

Complete the following tasks:

1. **Filter** the dataframe to Store 1 sales and **aggregate** over departments to compute the total sales per store
2. Plot the **rolling\_mean** for **Weekly\_Sales**
  - What general trends do you observe?
3. Compute the **1, 2, 52 autocorrelations** for **Weekly\_Sales** and/or create an autocorrelation plot
4. What does the autocorrelation plot say about the type of model you want to build?



EXERCISE

# ACTIVITY: WALMART SALES DATA

---

## DIRECTIONS



### EXERCISE

5. Split the weekly sales data in a training and test set - using 75% of the data for training
6. Create an AR(1) model on the training data and compute the mean absolute error of the predictions.
7. Plot the residuals - where are their significant errors?
8. Compute and AR(2) model and an ARMA(2, 2) model - does this improve your mean absolute error on the held out set?
9. Finally, compute an ARIMA model to improve your prediction error - iterate on the p, q, and parameters comparing the model's performance



---

**CONCLUSION**

---

# TOPIC REVIEW

---

## CONCLUSION

---

- Time-series models use previous values to predict future values, also known as forecasting
- AR and MA model are simple models on previous values or previous errors respectively
- ARMA combines these two types of models to account for both gradual shifts (due to AR models) and abrupt changes (MA models)

---

# CONCLUSION

---

- ARIMA models train ARMA models on differenced data to account for non-stationary data
- Note that none of these models may perform well for data that has more random variation
- For example, for something like iphone sales (or searches) which may be sporadic, with short periods of increases, these models may not work well

**COURSE**

---

**BEFORE NEXT CLASS**

---

## **BEFORE NEXT CLASS**

---

# **DUE DATE**

- Final Project, Part 4 due: Thursday (5/17)

---

## LESSON

---

Q & A

## **LESSON**

---

# **EXIT TICKET**

**DON'T FORGET TO FILL OUT YOUR EXIT TICKET**